



The OpenGL Rendering Pipeline





The Rendering Pipeline

- The process to generate two-dimensional images from given virtual cameras and 3D objects
- The pipeline stages implement various core graphics rendering algorithms
- Why should you know the pipeline?
 - Understand various graphics algorithms
 - Program low level graphics systems
 - Necessary for programming GPUs
 - Help analyze the performance bottleneck



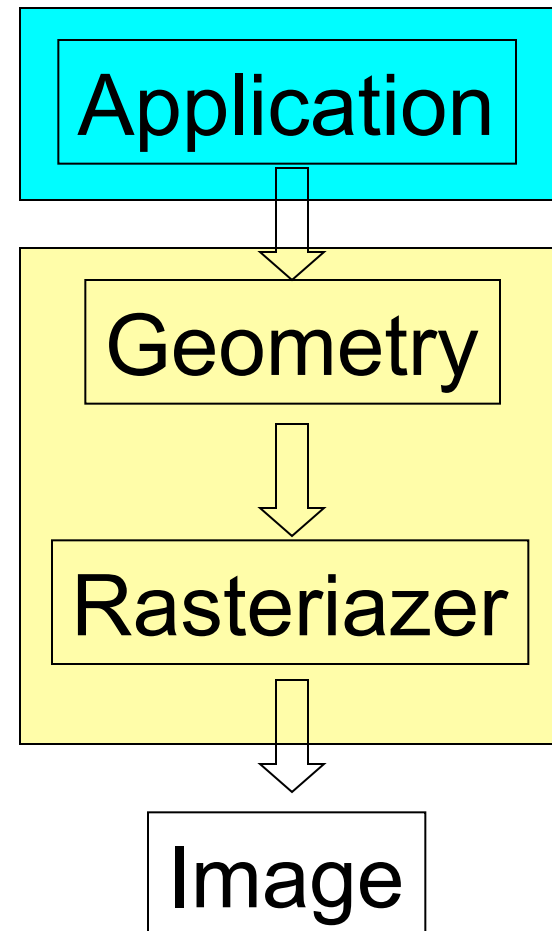

The Rendering Pipeline

- The basic construction – three conceptual stages
- Each stage is a pipeline and runs in parallel
- Graphics performance is determined by the slowest stage
- Modern graphics systems:

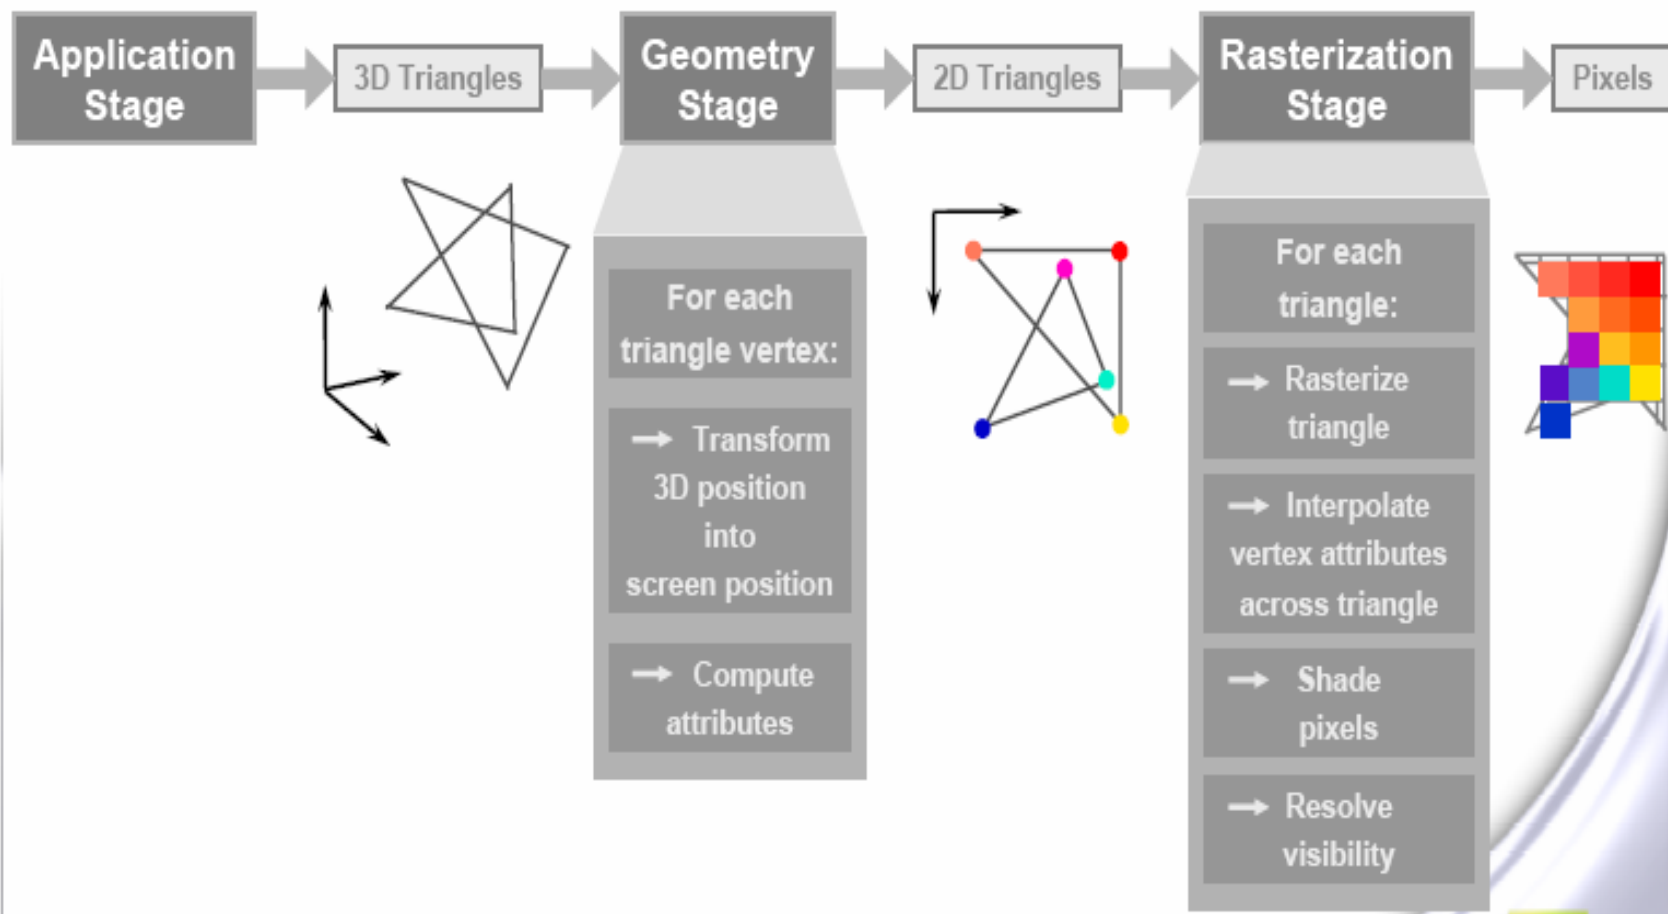

software:





hardware:



The Rendering Pipeline

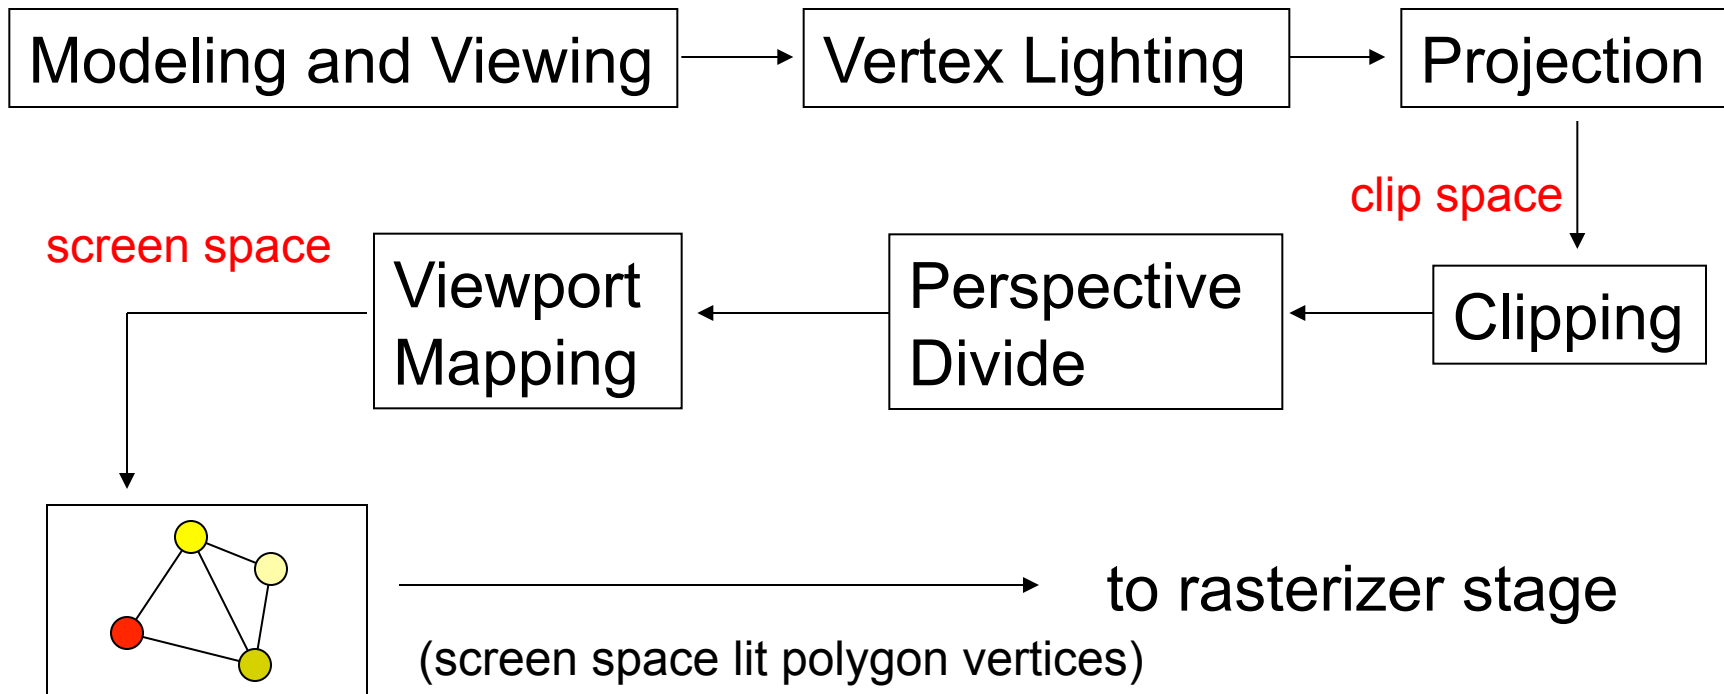


The Geometry Stage



(local space polygons)

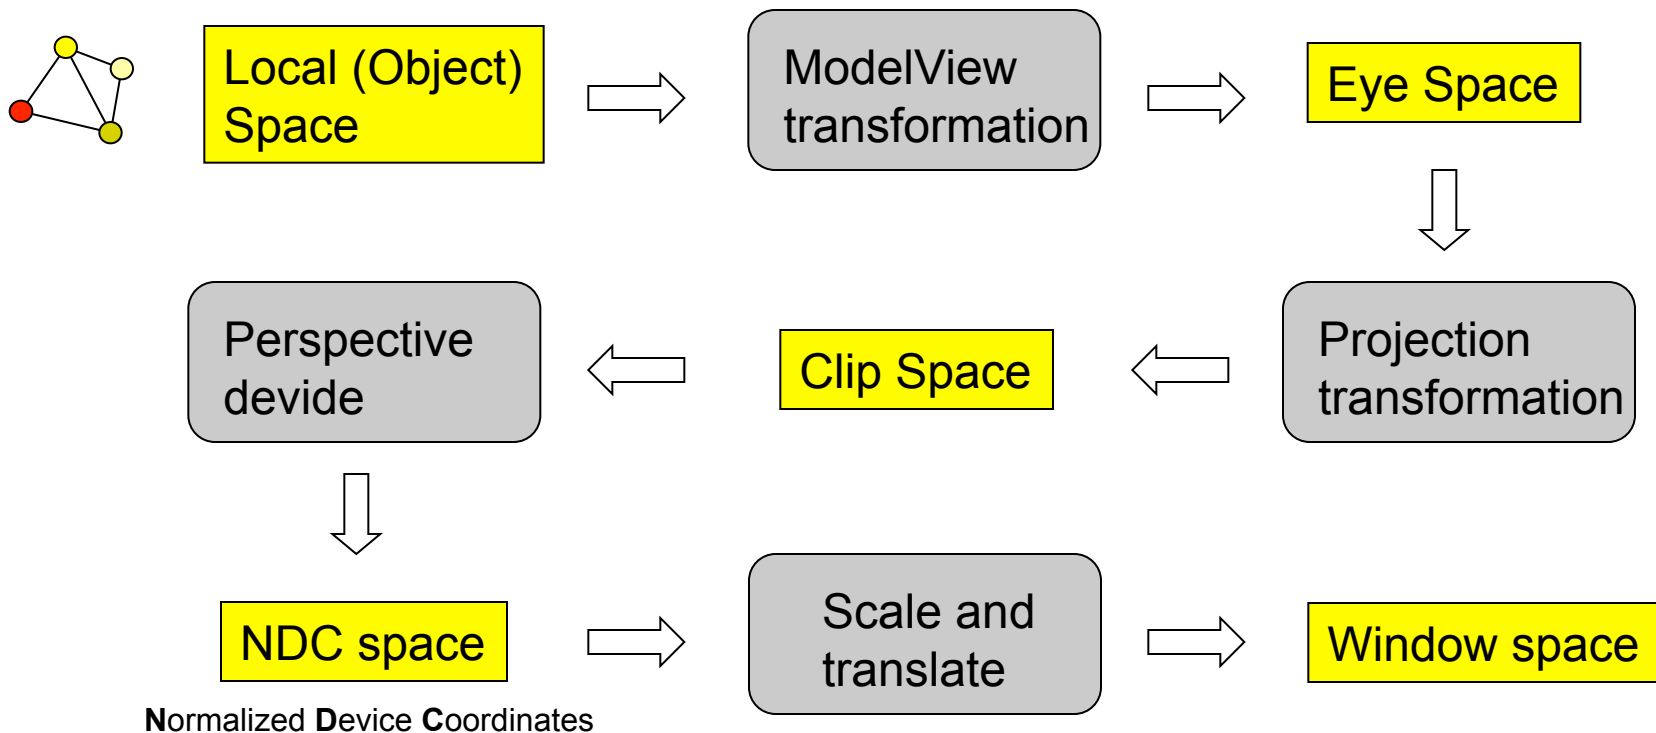
eye space






Transformation Pipeline

- Another view of the graphics pipeline






Different Spaces

- Local space
 - A space where you define the vertex coordinates, normals, etc. This is before any transformations are taking place
 - These coordinates/normals are multiplied by the OpenGL modelview (VM) matrix into the eye space
 - Modelview matrix: Viewing transformation matrix (V) multiplied by modeling transformation matrix (M), i.e., $GL_MODELVIEW = V * M$
 - OpenGL matrix stack is used to allow different modelview matrices for different objects




Different Spaces (cont'd)

- Eye space
 - Where per vertex lighting calculation is occurred
 - Camera is at $(0,0,0)$ and view's up direction is by default $(0,1,0)$
 - Light position is stored in this space after being multiplied by the OpenGL modelview matrix
 - Vertex normals are consumed by the pipeline in this space by the lighting equation



Different Spaces (cont'd)

- Clip Space
 - After projection and before perspective divide
 - Clipping against view frustum done in this space
 - $-W \leq X \leq W; -W \leq Y \leq W; -W \leq Z \leq W;$
 - New vertices are generated as a result of clipping
 - The view frustum after transformation is a parallelepiped regardless of orthographic or perspective projection
- Perspective Divide
 - Transform clip space into NDC space
 - Divide (x,y,z,w) by w where $w = z/-d$ ($d=1$ in OpenGL so $w = -z$)
 - Result in foreshortening effect



Different Spaces (cont'd)

- Window Space
 - Map the NDC coordinates into the window
 - X and Y are integers, relative to the lower left corner of the window
 - Z are scaled and biased to $[0,1]$
 - Rasterization is performed in this space
 - The geometry processing ends in this space



The Geometry Stage

- Transform coordinates and normal
 - Model->world
 - World->eye
- Normalize the normal vectors
- Compute vertex lighting
- Generate (if necessary) and transform texture coordinates
- Transform to clip space (by projection)
- Assemble vertices into primitives
- Clip against viewing frustum
- Divide by w (perspective divide if applies)
- Viewport transformation
- Back face culling




Introduce vertex
dependences ☹



The Rasterizer Stage

- Per-pixel operation: assign colors to the pixels in the frame buffer (a.k.a **scan conversion**)

- Main steps:




(frame buffer)

- Setup
- Sampling (convert a primitive to fragments)
- Texture lookup and Interpolation (lighting, texturing, z values, etc)
- Color combinations (illumination and texture colors)
- Fogging
- Other pixel tests (scissor, alpha, stencil tests etc)
- Visibility (depth test)
- Blending/compositing/Logic op



The Rasterization Stage

- Convert each primitive into fragments (not pixels)
- Fragment: transient data structures
 - position (x,y); depth; color; texture coordinates; etc




- Fragments from the rasterized polygons are then selected (z buffer comparison for instance) to form the frame buffer pixels




The Rasterization Stage

- Two main operations
 - Fragment selection: generate one fragment for each pixel that is intersected by the primitive
 - Fragment assignment: sample the primitive properties (colors, depths, etc) for each fragment - nearest neighbor continuity, linear interpolation, etc



Polygon Scan Conversion

- The goal is to compute the scanline-primitive intersections
- OpenGL Spec does not specify any particular algorithm to use
- Brute Force: try to intersect each scanline with all edges as we go from y_{min} to y_{max}
- We can do better
 - Find y_{min} and y_{max} for each edge and only test the edge with scanlines in between
 - For each edge, only calculate the intersection with the y_{min} ; calculate dx/dy ; calculate the new intersection as $y=y+1$, $x=x+dx/dy$
 - Change $x=x+dx/dy$ to integer arithmetic (such as using Bresenham's algorithm)



Rasterization steps

- Texture interpolation
- Color interpolation
- Fog (blend the fog color with the fragment color based on the depth value)
- Scissor test (test against a rectangular region)
- Alpha test (compare with alpha, keep or drop it)
- Stencil test(mask the fragment depending on the content of the stencil buffer)
- Depth test (z buffer algorithm)
- Alpha blending
- Dithering (make the color look better for low res display mode)